

COL774: Assignment 3

Sayam Sethi

October 2021

Contents

1	Random Forests	1
1.1	Grid Search	1
1.2	Paramter Sensitivity Analysis	2
1.2.1	n_estimators	2
1.2.2	max_features	3
1.2.3	min_samples_split	3
2	Neural Networks	5
2.1	Sigmoid Activation using a Fixed Learning Rate	5
2.1.1	5 Hidden Units	5
2.1.2	10 Hidden Units	6
2.1.3	15 Hidden Units	6
2.1.4	20 Hidden Units	7
2.1.5	25 Hidden Units	7
2.2	Sigmoid Activation using an Adaptive Learning Rate	7
2.2.1	5 Hidden Units	8
2.2.2	10 Hidden Units	8
2.2.3	15 Hidden Units	9
2.2.4	20 Hidden Units	9
2.2.5	25 Hidden Units	10
2.3	Multiple Hidden Layers: Comparing ReLU and Sigmoid	10
2.3.1	ReLU	10
2.3.2	Sigmoid	11
2.4	MLPClassifier (Scikit-Learn Library)	11
2.4.1	ReLU	11
2.4.2	Sigmoid	12

1 Random Forests

1.1 Grid Search

The best model obtained is for the following parameters:

- **n_estimators:** 450

- `max_features`: 0.3
- `min_samples_split`: 6

1.2 Paramter Sensitivity Analysis

1.2.1 `n_estimators`




Figure 1: Plot of the effect of `n_estimators` on the accuracy of the model

The validation accuracy increases on increasing the number of estimators. However the test accuracy drops slightly. However the training accuracy increases with increasing the number of estimators. Therefore, the model overfits with larger number of estimators.

1.2.2 max_features




Figure 2: Plot of the effect of max_features on the accuracy of the model

The accuracy decreases with the increase in max_features. However, for small increase, the accuracy improves slightly. Therefore, smaller values of max_features are preferred. The training accuracy increases for larger values of max_features. Therefore, larger values overfit the model.

1.2.3 min_samples_split

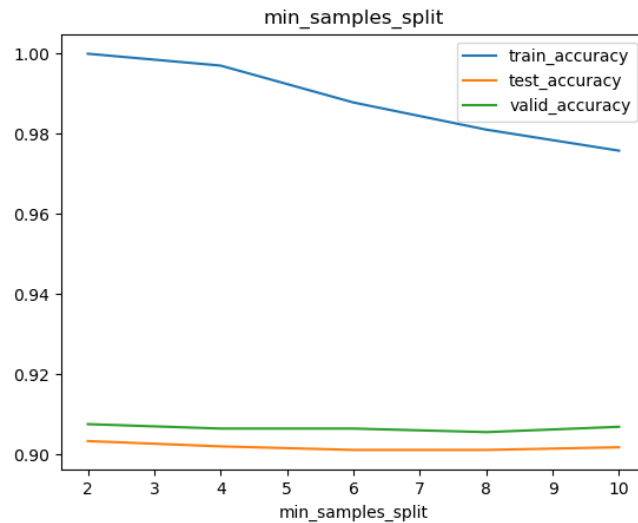


Figure 3: Plot of the effect of min_samples_split on the accuracy of the model

The training accuracy decreases with the increase in min_samples_split. However, the validation and testing accuracy increases with the increase in min_samples_split. Therefore, the model overfits

with smaller values of `min.samples.split`.


2 Neural Networks

2.1 Sigmoid Activation using a Fixed Learning Rate


The stopping criteria used is to reach a change in the average cost function across two epochs less than 10^{-8} or 1000 epochs, whichever happens earlier. When learning, it was observed that for the more complicated neural networks, the epoch criteria was met earlier.

The plots and confusion matrices are shown below. It is observed that the model improves when using larger number of perceptrons, however, it slightly drops on using 25 perceptrons. This is because 1000 epochs is too less to converge.

Additionally, we notice that no model learns the classes 2 – 9 at all. This is inferred from the zeroes in all the columns other than the first two in the confusion matrix.



(a) Plot of the training and test accuracy



(b) Plot of the training time

2.1.1 5 Hidden Units

training accuracy = 0.5723310675729708

test accuracy = 0.548671

$$\text{confusion matrix} = \begin{pmatrix} 377431 & 123778 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 251258 & 171240 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 21594 & 26028 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 7397 & 13724 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 2007 & 1878 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1746 & 250 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 352 & 1072 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 30 & 200 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 9 & 3 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 3 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{pmatrix} \quad (1)$$

training time = 39.02067995071411

2.1.2 10 Hidden Units

training accuracy = 0.6628548580567772

test accuracy = 0.625259

$$\text{confusion matrix} = \begin{pmatrix} 400963 & 100246 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 198202 & 224296 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 11667 & 35955 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 4343 & 16778 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 2129 & 1756 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1747 & 249 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 141 & 1283 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 6 & 224 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 10 & 2 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 2 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{pmatrix} \quad (2)$$

training time = 53.303688526153564

2.1.3 15 Hidden Units

training accuracy = 0.6614154338264694

test accuracy = 0.629156

$$\text{confusion matrix} = \begin{pmatrix} 395803 & 105406 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 189145 & 233353 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 11336 & 36286 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 2809 & 18312 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1754 & 2131 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1806 & 190 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 116 & 1308 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 229 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 6 & 6 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 2 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{pmatrix} \quad (3)$$

training time = 59.76480436325073

2.1.4 20 Hidden Units

training accuracy = 0.7426229508196721

test accuracy = 0.703346

$$\text{confusion matrix} = \begin{pmatrix} 416637 & 84572 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 135789 & 286709 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 3786 & 43836 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 942 & 20179 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 2269 & 1616 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1787 & 209 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 1423 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 230 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 6 & 6 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 2 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{pmatrix} \quad (4)$$

training time = 63.88574767112732

2.1.5 25 Hidden Units

training accuracy = 0.7303878448620552


test accuracy = 0.695839

$$\text{confusion matrix} = \begin{pmatrix} 435371 & 65838 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 162030 & 260468 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 6315 & 41307 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 2124 & 18997 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 2619 & 1266 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1857 & 139 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 20 & 1404 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 2 & 228 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 11 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 3 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{pmatrix} \quad (5)$$


training time = 71.55454468727112

2.2 Sigmoid Activation using an Adaptive Learning Rate

The stopping criteria used is the same as the previous section. The plots and confusion matrices are given below. It is observed that the accuracy is higher for 5 units of the hidden layer. However, the accuracy is smaller for the other cases. Additionally, the accuracy is lesser for 20 and 25 hidden units compared to 10, 15 hidden units. This is because using the adaptive learning rate slows down the learning and therefore 1000 epochs is not enough to converge. The training time increases as expected.



(a) Plot of the training and test accuracy



(b) Plot of the training time

2.2.1 5 Hidden Units

training accuracy = 0.589484206317473

test accuracy = 0.572393

$$\text{confusion matrix} = \begin{pmatrix} 387031 & 114178 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 237136 & 185362 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 19466 & 28156 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 5737 & 15384 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1329 & 2556 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1678 & 318 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 251 & 1173 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 18 & 212 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 4 & 8 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 2 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{pmatrix} \quad (6)$$

training time = 40.8355929851532

2.2.2 10 Hidden Units

training accuracy = 0.6117153138744502

test accuracy = 0.588626

$$\text{confusion matrix} = \begin{pmatrix} 381880 & 119329 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 215752 & 206746 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 16304 & 31318 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 4216 & 16905 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 2739 & 1146 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1761 & 235 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 210 & 1214 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 9 & 221 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 11 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 3 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{pmatrix} \quad (7)$$

training time = 58.18630528450012

2.2.3 15 Hidden Units

training accuracy = 0.6239904038384646

test accuracy = 0.601769

$$\text{confusion matrix} = \begin{pmatrix} 385665 & 115544 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 206394 & 216104 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 14607 & 33015 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 4136 & 16985 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 2545 & 1340 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1792 & 204 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 180 & 1244 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 10 & 220 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 10 & 2 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 2 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{pmatrix} \quad (8)$$

training time = 66.1306664943695

2.2.4 20 Hidden Units

training accuracy = 0.5859256297481008

test accuracy = 0.567364

$$\text{confusion matrix} = \begin{pmatrix} 392936 & 108273 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 248070 & 174428 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 19182 & 28440 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 6580 & 14541 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 2806 & 1079 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1809 & 187 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 273 & 1151 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 22 & 208 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 10 & 2 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 3 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{pmatrix} \quad (9)$$

training time = 68.35846948623657

2.2.5 25 Hidden Units

training accuracy = 0.6100359856057577

test accuracy = 0.589734

$$\text{confusion matrix} = \begin{pmatrix} 410562 & 90647 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 243326 & 179172 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 19587 & 28035 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 4789 & 16332 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 3881 & 4 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1760 & 236 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 291 & 1133 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 9 & 221 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 12 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 3 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{pmatrix} \quad (10)$$

training time = 81.43693113327026

2.3 Multiple Hidden Layers: Comparing ReLU and Sigmoid

The same stopping criteria is used for this part as well.

2.3.1 ReLU

training accuracy = 0.9233106757297082

test accuracy = 0.922804

$$\text{confusion matrix} = \begin{pmatrix} 500667 & 542 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 361 & 422137 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 47622 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 21121 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 3721 & 164 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1984 & 12 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1424 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 230 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 10 & 2 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 2 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{pmatrix} \quad (11)$$

training time = 195.58691668510437

Even though the model still doesn't learn classes 2 – 9, it does learn the classes 0 and 1 highly accurately. In the given training data, about 92% of the training data is in class 0 and 1 and the training accuracy is larger than 92%, most of which has contribution from these two classes. Therefore, ReLU is a very good choice for the given problem.

2.3.2 Sigmoid

training accuracy = 0.5003998400639744

test accuracy = 0.4974

$$\text{confusion matrix} = \begin{pmatrix} 482380 & 18829 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 407478 & 15020 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 46043 & 1579 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 20476 & 645 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 3707 & 178 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1879 & 117 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1389 & 35 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 225 & 5 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 12 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 3 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{pmatrix} \quad (12)$$

training time = 77.64794874191284

The accuracy drops to 50% in this case. The stopping criteria reached in this case was of change in cost function becoming smaller than 10^{-8} . Thus, we can conclude that the sigmoid derivative is not a very good choice since the backpropagation is smaller.

2.4 MLPClassifier (Scikit-Learn Library)

Models were trained for both ReLU and sigmoid activation functions.

2.4.1 ReLU

training accuracy = 0.9999200319872051

test accuracy = 0.989527

$$\text{confusion matrix} = \begin{pmatrix} 501126 & 5 & 0 & 0 & 75 & 3 & 0 & 0 & 0 & 0 \\ 107 & 422383 & 8 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 565 & 706 & 46328 & 23 & 0 & 0 & 0 & 0 & 0 & 0 \\ 924 & 6 & 163 & 20028 & 0 & 0 & 0 & 0 & 0 & 0 \\ 3694 & 1 & 0 & 0 & 190 & 0 & 0 & 0 & 0 & 0 \\ 1929 & 0 & 0 & 0 & 0 & 67 & 0 & 0 & 0 & 0 \\ 442 & 0 & 94 & 365 & 0 & 0 & 521 & 2 & 0 & 0 \\ 0 & 0 & 0 & 228 & 0 & 0 & 2 & 0 & 0 & 0 \\ 8 & 0 & 0 & 0 & 1 & 3 & 0 & 0 & 0 & 0 \\ 2 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \end{pmatrix} \quad (13)$$

training time = 9.497213125228882

The model almost perfectly learns the training data and performs very well on the test data. The training also happens almost instantly and many folds faster than our training (our model was in fact incomplete since the criteria of 1000 epochs was reached).

2.4.2 Sigmoid

training accuracy = 0.11095561775289885

test accuracy = 0.102564

$$\text{confusion matrix} = \begin{pmatrix} 378353 & 122856 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 320188 & 102310 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 36359 & 11263 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 16278 & 4843 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 2865 & 1020 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1384 & 612 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1096 & 328 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 172 & 58 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 9 & 3 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 2 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{pmatrix} \quad (14)$$

training time = 11.053844213485718

The accuracy is terrible in this case and much worse than what is learnt by our model. This further proves our claim that ReLU is a much better activation function for the given problem.